

Distributed Global Scheduling in Datacenters

Smita Vijayakumar

First Year Ph.D. Student

sv440@cst.cam.ac.uk

Evangelia Kalyvianaki

Ph.D. Supervisor

ek264@cst.cam.ac.uk

Anil Madhavapeddy

Ph.D. Supervisor

avsm2@cst.cam.ac.uk

State-of-the-Art Scheduling

Underutilised Datacenter Resources

Azure¹

- ❖ 60% VMs have $\leq 20\%$ CPU usage!

Alibaba²

- ❖ Average server CPU 50%
- ❖ Memory $\leq 60\%$

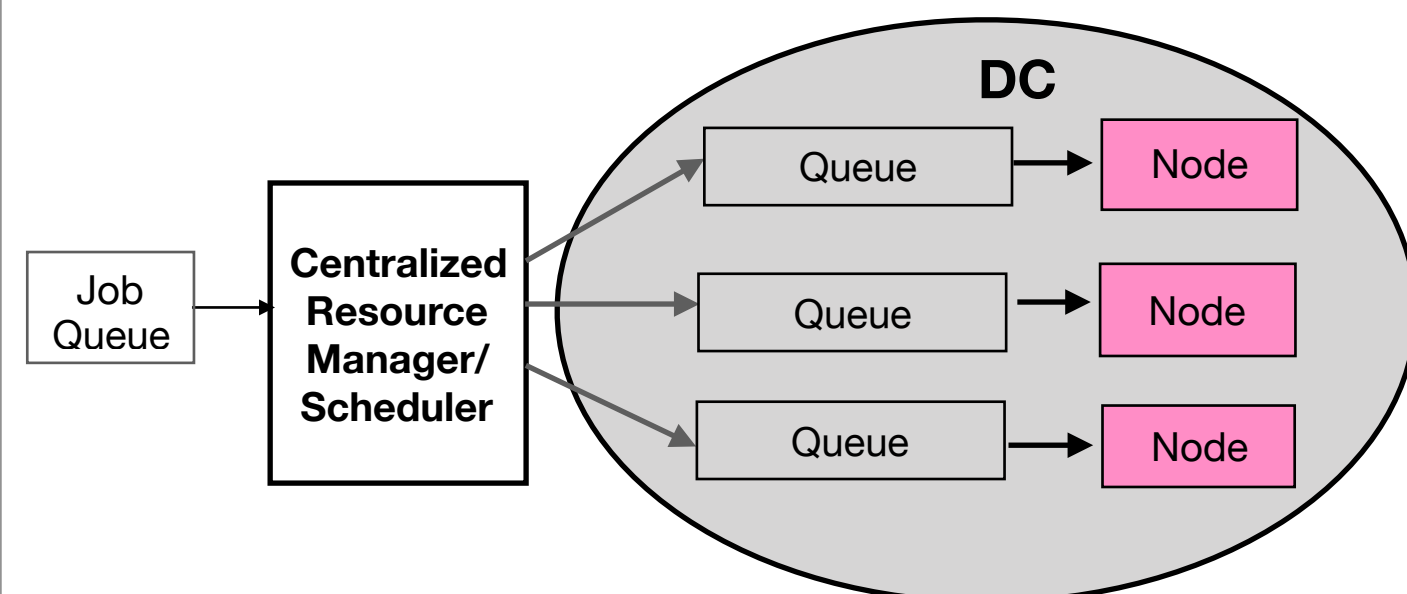
Underutilisation is expensive!³

¹[Resource Central, SOSP,'17]

²[https://github.com/alibaba/clusterdata]

³[Scalable system scheduling for HPC and big data, JPDC,17]

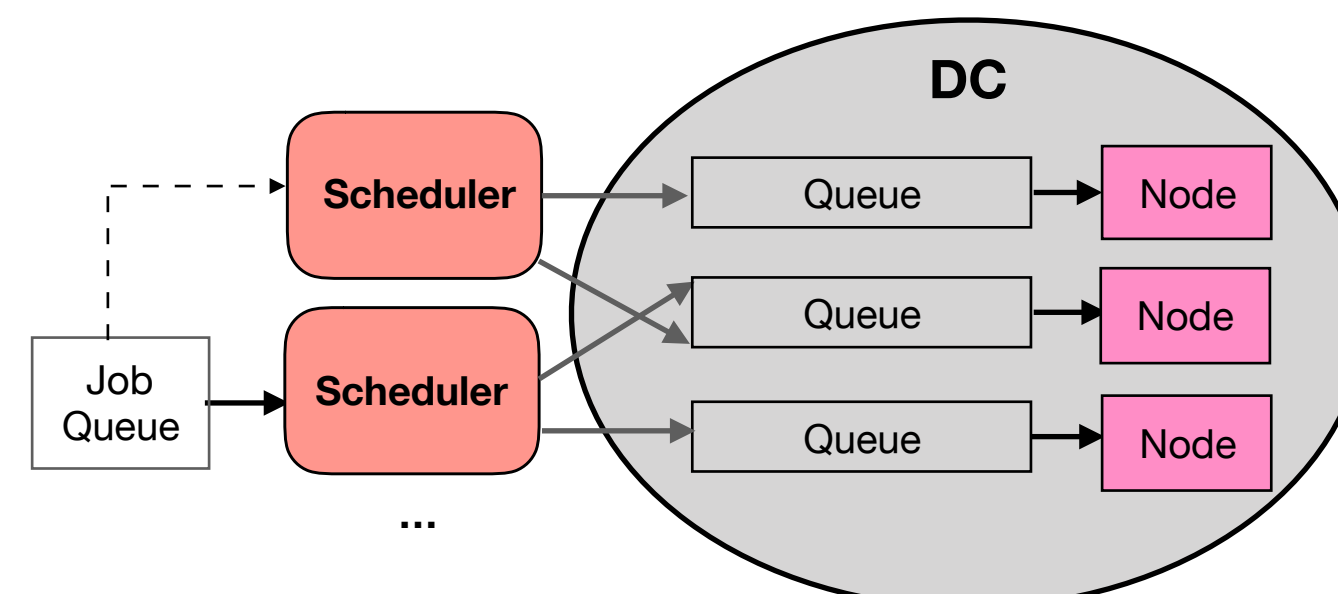
Centralised



Examples - Mesos [NSDI'11], Yarn, Apollo [OSDI'14]

- ✓ Global resource view
- ✗ Scheduler can be a bottleneck
- ✗ Delayed, high volumes of resource updates

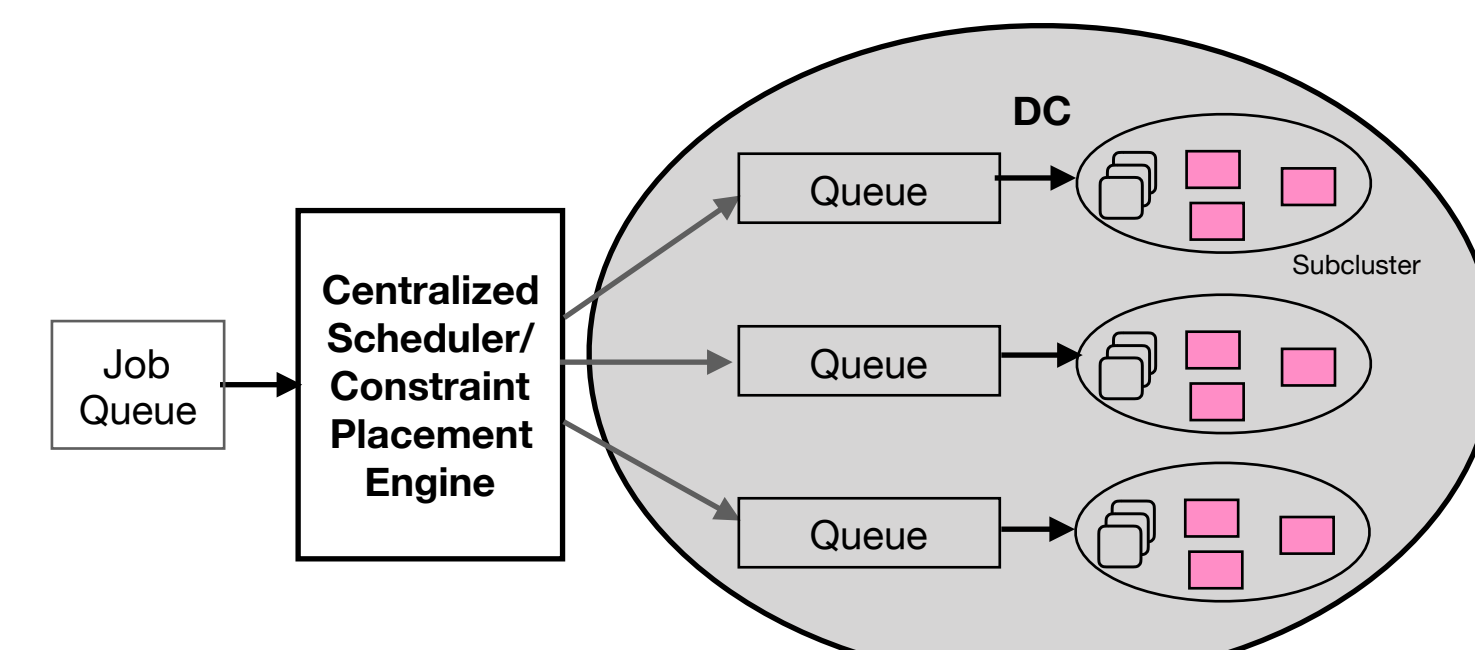
Decentralised



Example - Sparrow [NSDI'14]

- ✓ Fast and simple
- ✗ Unsuitable for long running jobs
- ✗ Not globally optimal

Hybrid

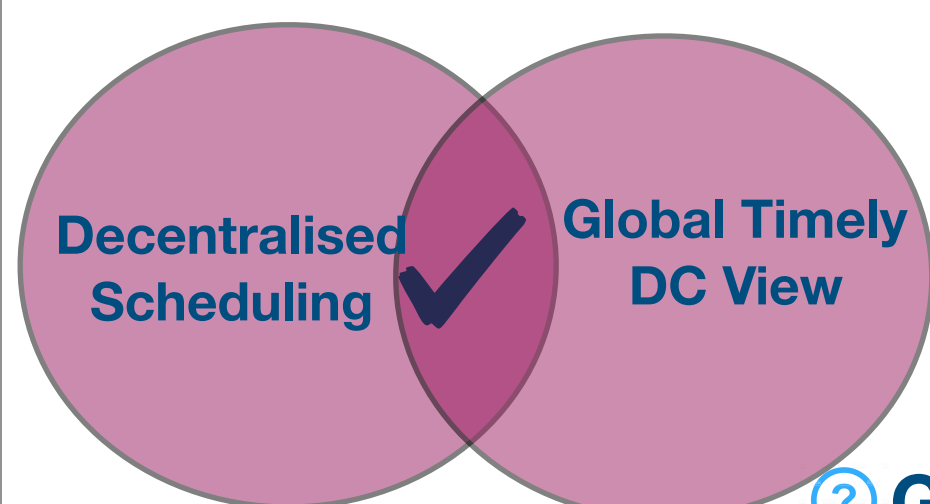


Example - Hydra [NSDI'19], Medea [EuroSys'18], Borg [EuroSys'15]

- ✓ Policy-driven job/task placement
- ✓ Less node information traffic
- ✗ Centralised or decentralised components

Proposed Direction

Global Scheduling at Node Level



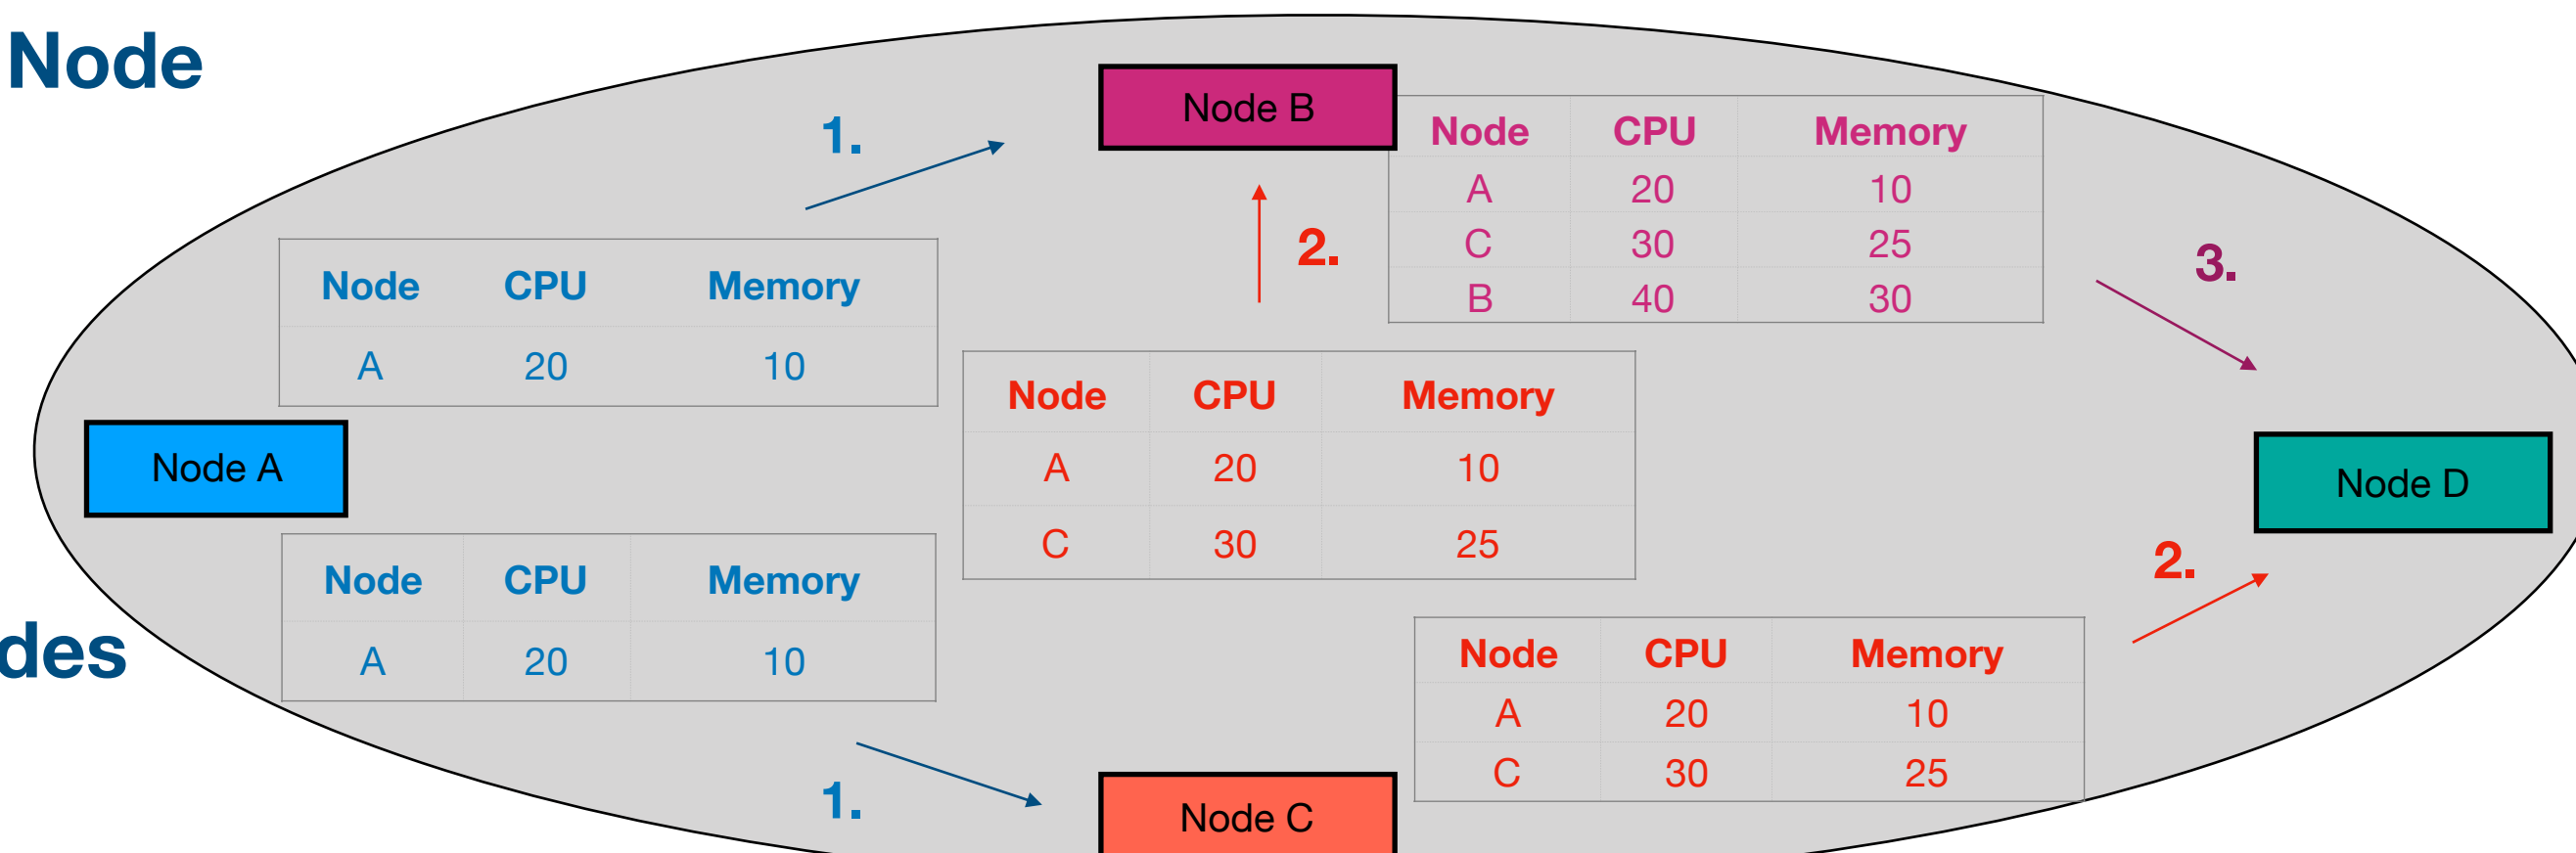
Challenges

- ? Good for long and short jobs
- ? Volume and frequency of updates
- ? Time from local to global view

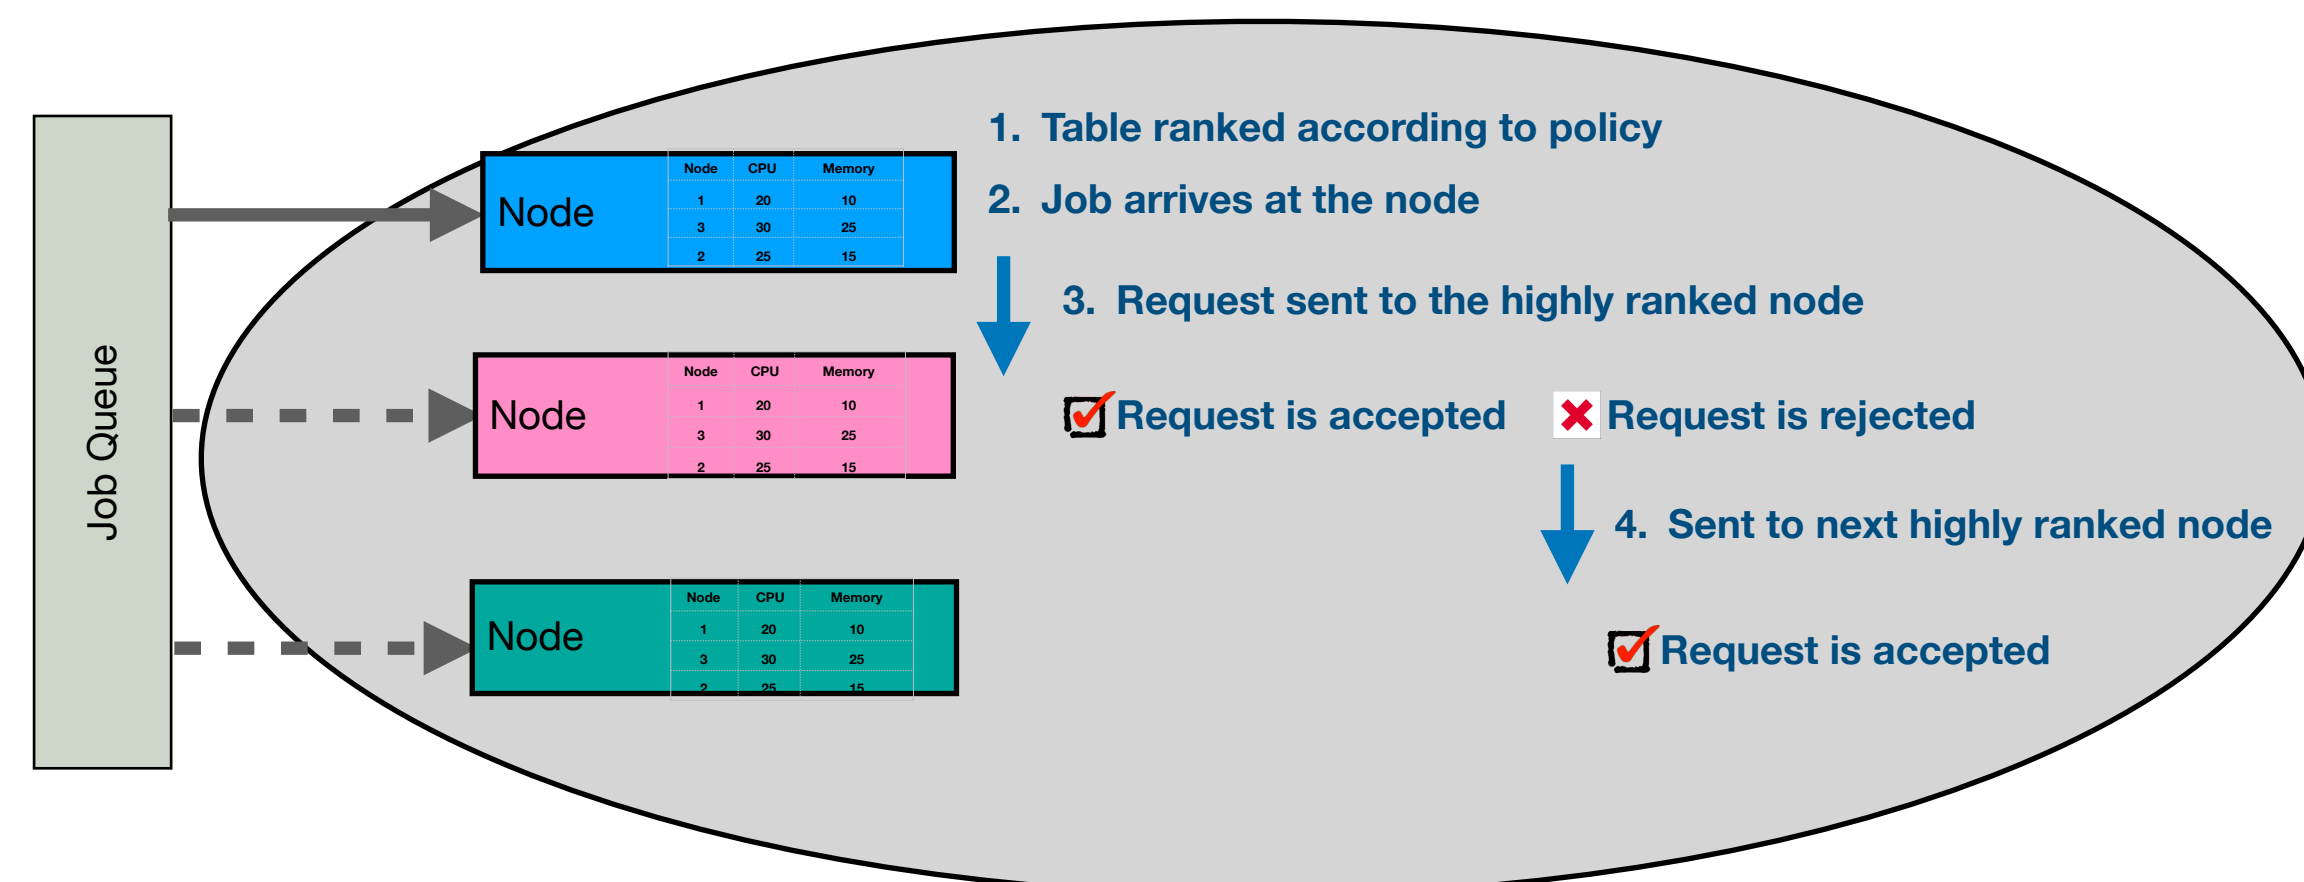
Up-to-Date Global View at Each Node

Inspired by routing protocols

- ✓ BGP, OSPF, ...
- ✓ Resource data propagation
- ✓ Global view convergence
- ✓ Same ranking policy across nodes



Intra-DC Scheduling



Challenges

- ? Collision avoidance
- ? Minimise inter-DC traffic
- ? Minimise scheduling time

Scheduling Using "Up-to-Date" Global View

Inter-DC Scheduling

